

OSVVM

Verification Component

Developer's Guide

Release 2022.06

By

Jim Lewis

SynthWorks VHDL Training

Jim@SynthWorks.com

<http://www.SynthWorks.com>

Table of Contents

| | | |
|------|--|----|
| 1 | Overview..... | 3 |
| 2 | Step 1: Identify the Transactions Supported by the Interface | 3 |
| 3 | Step 2: Map the VC transactions to OSVVM MIT | 4 |
| 4 | Step 3: Create the Verification Component | 4 |
| 5 | Create a simple blocking VC | 6 |
| 5.1 | Package References..... | 6 |
| 5.2 | DpRamController Entity Interface..... | 7 |
| 5.3 | Declare one or more AlertLogIDs..... | 8 |
| 5.4 | Anatomy of the TransactionHandler..... | 8 |
| 5.5 | Creating the Write Operation | 10 |
| 5.6 | Creating the Read Operation | 11 |
| 5.7 | Creating Directive Transactions | 11 |
| 5.8 | Detecting Multiple Drivers | 12 |
| 5.9 | Catching Unused Transactions..... | 12 |
| 5.10 | The Final Design | 12 |
| 6 | OSVVM Blocking VC vs a "Lite" Approach | 12 |
| 7 | Supporting Asynchronous Transactions | 13 |
| 8 | Running the Examples | 13 |
| 9 | About the OSVVM | 13 |
| 10 | About the Author - Jim Lewis | 14 |
| 11 | References..... | 14 |

1 Overview

This guide explores writing a verification component (VC) for a DpRam controller interface. The signaling on the DpRam controller is quite simple. This allows the focus to be on the aspects of creating an OSVVM style VC.

Be sure to read the OSVVM's Structured Testbench Framework Guide first. This guide provides you with the background to the overall structure of OSVVM's testbench and where verification components fit in.

2 Step 1: Identify the Transactions Supported by the Interface

Writing an OSVVM VC starts by looking at the interface we want to drive. In this guide, we will test the OSVVM DpRam block. This DpRam is similar to the memory models provided by FPGAs – such as Xilinx's BRAM.

The DpRam has two identical interfaces (hence Dual Port) that can both write or read from the interface. A block diagram for the DpRam is shown in Figure 1. It has registers on all input signals (AddrX, DataX, and WriteX) and optionally has a register on the output DataOutX.




Figure 1. DpRam structure

A timing diagram for the interface is shown in Figure 2. A write operation starts by presenting the Address (AddrX), Data (DataInX), and the Write indicator (WriteX). On the next rising edge of clock, the write operation is completed. A read operation starts by presenting Address (AddrX). On the next rising edge of clock the Address is accepted and read data (DataOutX) is available.

It has registers on all input signals (AddrX, DataX, and WriteX) and optionally has a register on the output DataOutX.




Figure 2. DpRam Timing Diagram

From a transaction point of view, OSVVM classifies interfaces that support address and data as an Address Bus interface. This interface supports a Write and a Read operation. Depending on the FPGA, it can also support a Write and Read of the same address. What you get on that read will depend on the FPGA target and configuration of the DPRAM for that target. For our purposes, the DPRAM provides the value that was in the RAM before the write operation started.

3 Step 2: Map the VC transactions to OSVVM MIT

To simplify VC development, OSVVM provides model independent transactions for Address Bus and Streaming interfaces that are intended to support a superset of the transactions that you need to implement in your VC.

We are in luck, the OSVVM Address Bus Model Independent Transactions support all the transactions needed in our interface. They are shown in Figure 3.

| |
|---|
| <code>Write(TRec, iAddr, iData, [MsgOn])</code> |
| <code>Read(TransactionRec, iAddr, oData, [StatusMsgOn])</code> |
| <code>ReadCheck(TransactionRec, iAddr, iData, [StatusMsgOn])</code> |
| <code>WriteAndRead(TransactionRec, iAddr, iData, oData, [StatusMsgOn])</code> |

Figure 3. Address Bus Model Independent Transactions (a subset)

4 Step 3: Create the Verification Component

We have choices with respect to the complexity of verification component that we write. The simplest VC we can write is a blocking verification component. In this case, when a process in the test sequencer calls a transaction, that transaction call blocks (does not return) until the VC is finished executing the transaction. Blocking VC are sufficient for interfaces that only do one transaction at a time – such as a UART or AxiStream.

A more complex VC is an asynchronous VC. In this case, the test sequencer sends a transaction to the VC, but does not wait for the VC to complete the transaction. This level of complexity is needed

support an interface that supports multiple independent operations at a time – such as an AXI4 interface where the Write Address, Write Data, Write Response, Read Address, and Read Data aspects of the interface are all handled independently.

The choice of blocking vs asynchronous VC also impacts the structure of the test sequencer.

In the test sequencer, when we call transactions, we can either use multiple processes (OSVVM's preferred method) or use a single process. Both have advantages and disadvantages. Independent of whether we use multiple processes or a single process, test cases must be able to create simultaneous actions on independent interfaces – for example have activity on two separate AXI4 interfaces start at the exact same clock cycle.

OSVVM prefers a test sequencer that uses a separate process (one for each independent VC – aka multiple processes) to call transactions as it allows the usage of simpler blocking VC while at the same time supporting the ability to create simultaneous actions on independent interfaces. The disadvantage is that to create a particular alignment of transactions, we need to synchronize (align) the independent processes. To accomplish this OSVVM has created synchronization primitives, such as `WaitForBarrier`, in the Utility library package `TbUtilPkg`.

For a test sequencer that uses a single process to call transactions, the VC must support asynchronous handling of the transactions in order to support simultaneous actions on interfaces. Generally single process dispatch is more complicated. Simple things that naturally require multiple clock cycles, like a read operation, must be refactored into an asynchronous read address dispatch followed by one or more try read data operations (repeated until the data is received). Going further when we start randomizing patterns on independent interfaces, it is much easier to do with multiple processes.

In the OSVVM Verification Component Library all AXI (Axi4 Full, Axi4 Lite and AxiStream) VC support both blocking and asynchronous transactions. For Axi4 this is a requirement of the interface. For all this means they can support either multiple process or single process transaction dispatch.

Part 1 shows the creation the creation of a simple blocking VC for the `DpRamController`.

5 Part 1: Create a simple blocking VC

If OSVVM were to have a "Lite" approach, this would be it. The structure for the DpRamController blocking VC is shown in see Figure 4. This VC is an address bus VC, and hence, uses OSVVM's Address Bus Model Independent Transactions.

```
entity DpRamController is
  port ( . . . ) ;
end entity DpRamController ;
architecture blocking of DpRamController is

  TransctionHandler : process
  begin

    wait for Transaction( . . . ) ;
    case TransRec.Operation is
      when WRITE_OP =>          -- Do Write things
      when READ_OP  =>          -- Do Read Things
      when ...    =>            -- Do Model Directives
    end case
  end process TransactionHandler ;
end architecture blocking ;
```

Figure 4. Blocking VC Structure

The remainder of this section creates the DpRamController VC piece by piece.

5.1 Package References

Figure 5 shows the package references for the DpRamController. The library osvvm_common and its context reference supports usage of OSVVM's Model Independent Transactions. The library osvvm and its context reference supports usage of the OSVVM Utility library. The IEEE library and its package references support general VHDL modeling.

```
library ieee ;
  use ieee.std_logic_1164.all ;
  use ieee.numeric_std.all ;
  use ieee.numeric_std_unsigned.all ;
  use ieee.math_real.all ;

library osvvm ;
  context osvvm.OsvvmContext ;

library osvvm_common ;
  context osvvm_common.OsvvmCommonContext ;
```

Figure 5. Package References

5.2 DpRamController Entity Interface

The DpRamController entity interface has the DUT interface and a transaction record interface. This is shown in Figure 6. The transaction record tells the VC what to do. For clocked interfaces, each output changes after a propagation delay after clock rises. There is one generic of the form `tpd_Clk_<output>` to specify a delay for each output.

```
entity DpRamController is
generic (
  MODEL_ID_NAME      : string := "" ;
  DEFAULT_DELAY      : time   := 1 ns ;
  tpd_Clk_Address    : time   := DEFAULT_DELAY ;
  tpd_Clk_Write      : time   := DEFAULT_DELAY ;
  tpd_Clk_oData      : time   := DEFAULT_DELAY
) ;
port (
  Clk          : In   std_logic ;
  nReset       : In   std_logic ;

  -- DpRam Functional Interface
  Address      : Out  std_logic_vector ;
  Write        : Out  std_logic ;
  oData        : Out  std_logic_vector ;
  iData        : In   std_logic_vector ;

  -- Address Bus Transaction Interface
  TransRec     : InOut AddressBusRecType
) ;

-- Name for OSVVM Alerts
constant MODEL_INSTANCE_NAME : string :=
  IfElse(MODEL_ID_NAME /= "",
    MODEL_ID_NAME, PathTail(to_lower(DpRamController'PATH_NAME))) ;

end entity DpRamController ;
```

Figure 6. DpRamController Entity Interface

5.3 Declare one or more AlertLogIDs

AlertLogIDs are used with self-checking, error handling, and message filtering (see AlertLogPkg User Guide for details). All OSVVM VC declare a signal ModelID of AlertLogIDType as the primary error handling and message filtering ID for this VC. Naming the signal ModelID is a naming convention and is not required – but does simplify copying code between VC if/when necessary. Additional IDs can be used to further identify the source of errors or messages within a VC.

Using a signal as shown in Figure 7 is the long way to do this. Many simulators allow you to declare ModelID as a constant and directly initialize it with the value from NewID. Unfortunately, the current version of the language does not permit this.

```
architecture SimpleBlocking of DpRamController is
  signal ModelID : AlertLogIDType ;
begin
  Initialize : process
  begin
    ModelID <= NewID(MODEL_INSTANCE_NAME) ;
    wait ;
  end process Initialize ;
```

Figure 7. Declaring and Initializing the AlertLogID

5.4 Anatomy of the TransactionHandler

The TransactionHandler process decodes commands from the test sequencer and executes them. The basic structure has a call to WaitForTransaction and then a case statement to decode the operation. In the body of the case target are interface actions that are specific to the interface. See Figure 8.

```
TransactionHandler : process
begin
  WaitForTransaction(
    Clk      => Clk,
    Rdy      => TransRec.Rdy,
    Ack      => TransRec.Ack
  ) ;
  Operation := TransRec.Operation ;
  case TransRec.Operation is
    when WRITE_OP =>      -- Do Write things
    when READ_OP  =>      -- Do Read Things
    when ...      =>      -- Do Model Directives
  end case ;
end process TransactionHandler ;
```

Figure 8. Basic Transaction Handler

The WaitForTransaction does the handshaking with the test sequencer and causes the process to wait until a transaction is available. It requires Clk and the two handshaking signals Rdy (driven by the sequencer) and Ack (driven by the VC). Both Rdy and Ack are in the transaction record. After checking Rdy and Ack, if the process is not aligned to clock, WaitForTransaction ensures that it is.

The case statement decodes the operation field that is in the transaction record.

If the process needs to initialize interface outputs, add an assignment to the outputs at the beginning of the process and put the repetitive part of the process into a "loop – end loop;". This is shown in Figure 9.

```

TransactionHandler : process
begin
  -- Initialize Outputs
  Address      <= (others => 'X') ;
  Write        <= 'X' ;
  oData        <= (others => 'X') ;

  loop
    wait for Transaction(. . . ) ;
    case TransRec.Operation is
      when WRITE_OP =>          -- Do Write things
      when READ_OP  =>          -- Do Read Things
      when ...    =>            -- Do Model Directives
    end case ;
  end loop ;
end process TransactionHandler ;

```

Figure 9. Transaction Handler with Initializations at the beginning.

5.5 Creating the Write Operation

The Write Operation implements the interface behavior for a write cycle. When WaitForTransaction exits, the VC will be aligned to the rising edge of clock. Address and oData will be driven aligned to clock with a propagation delay specified by the output timing generics (tpd_Clk_<output>). See Figure 10.

The Address for the write operation is in the Address field of the transaction record. SafeResize is used to optionally resize the value as well as convert from the type used in the record (osvvm.ResolutionPkg.std_logic_vector_max) to std_logic_vector. Write data is in the DataToModel field of the transaction record.

```
when WRITE_OP =>
  Address <= SafeResize(TransRec.Address, Address'length) after tpd_Clk_Address ;
  oData    <= SafeResize(TransRec.DataToModel, oData'length) after tpd_Clk_oData ;
  Write    <= '1' after tpd_Clk_Write ;

  WaitForClock(Clk) ;
  Log( . . . ) ; -- shown below
  Address <= not Address after tpd_Clk_Address ;
  oData    <= not oData    after tpd_Clk_oData ;
  Write    <= '0' after tpd_Clk_Write ;
```

Figure 10. Write Transaction Implementation

In a "Lite" based approach that implements interface behavior in procedures, the above code is the same code that would be in the procedure. Hence, writing a simple blocking VC is no more complex than writing a subprogram.

Logs are messages that provide feedback about what is happening in the VC. Logs are different from ordinary printing in that they can be filtered – meaning enabled or disabled by the test sequencer. Logs should be placed where they give information about what the VC is currently doing – such as this one which is at the point where the write operation is acted upon by the DpRam. The details of the log are shown in Figure 11. For details on Alert, Affirm, and Log please see the AlertLogPkg Users Guide. The final parameter to Log below is the transaction enable for the log. This allows the Log to be enabled for a single transaction without having to enable and then disable it.

```
Log( ModelID,
  "Write Operation, Address: " & to_hxstring(Address) &
  " Data: " & to_hxstring(oData) &
  " Operation# " & to_string (TransRec.Rdy)
  INFO,
  TransRec.StatusMsgOn
) ;
```

Figure 11. Details of the Call to Log.

5.6 Creating the Read Operation

The read operation, shown Figure 12, is similar to the write operation except we have to wait one to two clock cycles for the data to be available. The current DpRamController only waits for one clock cycle (and hence ignores the extra optional register on the output of the DpRam). Read data is returned to the test sequencer in the DataFromModel field of the transaction record. This version of SafeResize optionally resizes the data value as well as converts from std_logic_vector to the type used in the record (osvvm.ResolutionPkg.std_logic_vector_max).

```

when READ_OP | READ_CHECK =>
  Address <= SafeResize(TransRec.Address, Address'length) after tpd_Clk_Address ;
  Write   <= '0' after tpd_Clk_Write ;

  WaitForClock(Clk) ;
  Address <= not Address after tpd_Clk_Address ;

  WaitForClock(Clk) ;

  TransRec.DataFromModel <= SafeResize(iData, TransRec.DataFromModel'length) ;
  Log( . . . ) ;

```

Figure 12. Read Transaction Implementation

The Log for a simple read transaction is the same as the Log for the write operation shown previously.

5.7 Creating Directive Transactions

A directive transaction communicates with the verification component but does not result in any actions on the DUT interface. Common directive transactions for OSVVM include, WaitForClock, GetAlertLogID, and WaitForTransaction. WaitForClock stops the VC for the specified number of clocks. GetAlertLogID looks up the AlertLogID of the VC and returns it. WaitForTransaction waits for any pending transactions to complete and returns – in a blocking model, this does nothing because there all transactions complete. Implementation of these transactions is shown in Figure 13.

```

when WAIT_FOR_CLOCK =>
  WaitForClock(Clk, TransRec.IntToModel) ;

when GET_ALERTLOG_ID =>
  TransRec.IntFromModel <= integer(ModelID) ;

when WAIT_FOR_TRANSACTION =>
  null ;

```

Figure 13. Implementation of Directive Transactions

5.8 Detecting Multiple Drivers

When writing tests, it is natural to copy a piece of code from one process to another. It is common to forget to change the transaction record name to the one for the current process. As a result, a transaction record can end up with two processes in the test sequencer driving it. When the resolution function used for the operation field detects more than one operation being driven, it resolves (changes) the operation to `MULTIPLE_DRIVER_DETECT`.

It is best practice for the model to decode the `MULTIPLE_DRIVER_DETECT` and produce an appropriate `FAILURE` message. The implementation of multiple driver detect is shown in Figure 14.

```
when MULTIPLE_DRIVER_DETECT =>
  Alert(ModelID, MODEL_NAME & ": Multiple Drivers on Transaction Record." &
    " Transaction # " & to_string(TransRec.Rdy), FAILURE) ;
```

Figure 14. Implementation of Multiple Driver Detect

5.9 Catching Unused Transactions

A VC is not required to implement all transactions that are supported by OSVVM Address Bus Model Independent Transactions. It is best practice to produce an appropriate `FAILURE` message when an unsupported transaction is received. This is shown in Figure 15.

```
when others =>
  Alert(ModelID, "Transaction " & to_string(Operation) &
    " is not implemented by the DpRamController VC", FAILURE) ;
```

Figure 15. Catching and Signaling Unused Transactions

5.10 The Final Design

The final design can be found in the file `DpRamController_Blocking.vhd` in the directory `OsvvmLibraries/DpRam/src`.

6 OSVVM Blocking VC vs a "Lite" Approach

Some VHDL "Lite" based approaches implement the interface behavior in procedures rather than using an entity – like OSVVM's blocking VC. The thought is that it is easier. With OSVVM we have reduced the effort in creating the verification approach so that the code in OSVVM's blocking VC is no more complex than writing a procedure.

With incremental learning and coding an OSVVM blocking VC can be transformed into a VC that supports asynchronous transactions. The blocking VC can even be used temporarily in the test case generation until the VC that supports asynchronous transactions is done.

7 Part 2: Supporting Asynchronous Transactions

Some interfaces, such as this one, allow multiple things to be going on at a time. For example, what happens when we do a read address 1 followed by a write address 2. The interface allows the write address 2 to start on the clock immediately after starting the read address 1 operation, such that due to latencies in the system, both will finish at the same time. With blocking transactions, the read address 1 is required to finish before the write address 2 can start. This leaves a dead cycle on the interface.

Writing a full example of this for the DpRamController is an activity for another day. For now, both the AxiStream and Axi4 FULL support this capability. AxiStream can be found in the files AxiStreamTransmitter.vhd and AxiStreamReceiver.vhd in the directory OsvvmLibraries/AXI4/AxiStream/src. Axi4 Full can be found in the directory OsvvmLibraries/AXI4/Axi4/src.

This information can also be found in SynthWorks' class, Advanced VHDL Testbenches and Verification.

8 Running the Examples

See the steps in the OSVVM Overview guide for running the OSVVM demos. Do the steps shown in Figure 16 in your simulator. StartUp.tcl needs to be sourced each time you start the simulator. See the Script User Guide for additional details for Aldec's ActiveHDL, GHDL, Synopsys VCS, and Cadence Xcelium.

```
cd sim
source ../OsvvmLibraries/Scripts/StartUp.tcl
build ../OsvvmLibraries/OsvvmLibraries.pro
build ../OsvvmLibraries/DpRam/RunAllTests.pro
```

Figure 16. Compiling and Running DPRAM tests

9 Summary

In OSVVM, a verification component is nothing more than behavioral code. The busy work steps of creating a transaction API and interface are taken care of by OSVVM's Model Independent. OSVVM's approach is similar enough to RTL that an engineer can easily do both RTL and verification tasks.

10 About the OSVVM

The OSVVM utility and verification component libraries were developed and are maintained by Jim Lewis of SynthWorks VHDL Training. These libraries evolved from methodology and packages developed for SynthWorks' VHDL Testbenches and verification class.

Please support OSVVM by purchasing your VHDL training from SynthWorks.

11 About the Author - Jim Lewis

Jim Lewis, the founder of SynthWorks, has thirty plus years of design, teaching, and problem solving experience. In addition to working as a Principal Trainer for SynthWorks, Mr Lewis has done ASIC and FPGA design, custom model development, and consulting.

Mr. Lewis is chair of the IEEE 1076 VHDL Working Group (VASG) and is the primary developer of the Open Source VHDL Verification Methodology (OSVVM.org) packages. Neither of these activities generate revenue. Please support our volunteer efforts by buying your VHDL training from SynthWorks.

If you find bugs these packages or would like to request enhancements, you can reach me at jim@synthworks.com.

12 References

[1] Jim Lewis, Advanced VHDL Testbenches and Verification, student manual for SynthWorks' class.